

OrgLang

Reference

- [OrgLang Reference](#)
 - [Introduction](#)
 - [Installation](#)
 - [Concepts](#)
 - [Tables](#)
 - [Resources](#)
 - [Hybrid Scheduler](#)
 - [Flux](#)
 - [Uniformity](#)
 - [Notation](#)
 - [Lexical analysis](#)
 - [Character encoding](#)
 - [Line structure](#)
 - [Identifiers and keywords](#)
 - [Literals](#)
 - [Operators](#)
 - [Delimiters](#)
 - [Data model](#)
 - [Values and types](#)
 - [The standard type hierarchy](#)
 - [Special names](#)
 - [Execution model](#)
 - [Naming and binding](#)
 - [Closures and Scoping](#)
 - [Errors](#)
 - [Arithmetic conversions](#)
 - [Atoms](#)
 - [Unary arithmetic and bitwise operations](#)
 - [Binary arithmetic operations](#)
 - [Shifting operations](#)
 - [Binary bitwise operations](#)
 - [Comparisons](#)
 - [Boolean operations](#)
 - [Access and Conditional Expressions](#)
 - [Lambdas \(Anonymous Operators\)](#)
 - [Table Construction](#)
 - [Evaluation Order](#)
 - [Operator Precedence](#)
 - [Assignment](#)
 - [Resources Execution](#)

- [Flux Operators](#)
- [Memory management](#)
- [Operator Definitions](#)
- [Module Definitions](#)
- [Full Grammar Specification](#)
- [Build Model](#)
- [Imports](#)
- [Project Structure](#)
- [Future Goals \(Wishlist\)](#)
- [Non-Goals \(v1\)](#)

OrgLang Reference

Introduction

OrgLang is a programming language designed with a focus on **orthogonality**, **aesthetics**, and **minimalism**. It is built on a few deeply integrated concepts:

1. **Uniformity**: Everything in OrgLang is data. The fundamental data structure is the **Table**—a collection of pairs that represents scopes, modules, and objects.
2. **Flux**: Execution is modeled as a flow of data through transformations. The `->` operator (Push) is as central as variable assignment.
3. **Reification**: Side effects and external resources (Files, Networks, System calls) are reified as **Resources**, allowing for deterministic and scoped management of effects.

OrgLang aims to provide a premium development experience where the code is as beautiful as the systems it builds.

Installation

OrgLang is currently in development and is not yet available for installation. Stay tuned for updates!

Concepts

Tables

Tables are the fundamental data structure in OrgLang. They are collections of pairs that represent scopes, modules, and objects.

Resources

Resources are used to represent side effects and external resources (Files, Networks, System calls) in OrgLang. They are reified as **Resources**, allowing for deterministic and scoped management of effects.

The execution of a program is a sequence of transformations that are applied to a **Table**. Each transformation is a function that takes a **Table** as input and returns a **Table** as output.

The typical `main` function of most languages in OrgLang would be:

```
main: {@args -> "Hello, World!" -> @stdout};
```

While a more robust example would be:

```
main: {@args -> "Hello, World!" -> @stdout ?? @stderr};
```

Every interaction with the outside world is a resource. The builtin `args` resource is the entry point of an executable. The `args` resource pulls data from the command line arguments and environment variables. A more complex example would be:

```
main: {@args -> "Hello, $0!" |> $ -> @stdout ?? @stderr};
```

The `|>` operator is the Partial Application operator. It is used to apply a function to a value and return a new function that can be applied to another value. This allows us to create a unary operator that uses the value pulled from the resource on the left - `args` - as an argument.

The `args` resource sends a Table with the command line arguments and environment variables to the partial function on the right, and the String substitution operator `$` takes its arguments from the Table.

Hybrid Scheduler

OrgLang uses a hybrid scheduler that combines **M:N Scheduling** (similar to Go's goroutines, but optimized for our Arena memory model) with **Completion-Based** execution. The scheduler treats every pulse of data as a task that can be suspended and resumed based on hardware availability.

So, the `@args` function is simply the “Seed Pulse” that starts the chain reaction. If we create two pulses, one for each core, the scheduler will execute them in parallel, and when one is finished, it will be automatically deallocated.

```
main: {@args -> ["Hello, World!" -> @stdout  
                  "Burn the World!" -> @stderr]};
```

This is valid code that will print “Hello, World!” to stdout and “Burn the World!” to stderr, with no guarantee of order.

Flux

Flux is the flow of data through transformations in OrgLang. The `->` operator (Push) is as central as variable assignment.

Uniformity

Everything in OrgLang is data. The fundamental data structure is the **Table**—a collection of pairs that represents scopes, modules, and objects.

Notation

This manual uses the following conventions:

- **Bold text** highlights key concepts and terms defined for the first time.
- **Monospace text** denotes identifiers, keywords, operator symbols, and literal values.
- Code examples are provided in blocks with syntax highlighting.
- Grammar specifications are written in **EBNF** (Extended Backus-Naur Form). > **[!NOTE]** > Alerts like this provide additional context, implementation details, or helpful tips.

Lexical analysis

Lexical analysis is the first stage of the OrgLang compiler. An OrgLang program is read as a sequence of characters, which are then grouped into meaningful units called **tokens**.

The lexer (or tokenizer) performs this transformation by scanning the source text from beginning to end. It identifies various types of tokens:

- **Identifiers and Keywords:** Symbolic names for variables and functions.
- **Literals:** Constant values such as numbers and strings.
- **Operators:** Symbols representing computations or flows.
- **Delimiters:** Structural symbols like parentheses and semicolons.

Character encoding

OrgLang source files are expected to be encoded in **UTF-8**. While the current implementation primarily focuses on the ASCII subset for structural elements and identifiers, UTF-8 support ensures that strings and comments can contain any Unicode character.

Line structure

Comments

OrgLang supports two types of comments: single-line and multiline (block) comments.

Single-line comments start with the hash character (#) and extend to the end of the line. They are ignored by the compiler.

```
# This is a single-line comment  
x : 42; # Comment after an expression
```

Multiline comments (also known as block comments) are enclosed in three consecutive hash characters (###).

[WARNING] The multiline comment marker ### must start at the first column of the line.

Everything between the opening and closing ### markers is treated as a comment and ignored.

```
###  
This is a multiline comment.  
It can span multiple lines.  
###
```

Blank lines

A line that contains only whitespace (spaces, tabs, and form feeds) is considered a blank line and is ignored by the compiler. Blank lines are recommended to separate logical blocks of code and improve readability.

Indentation

Unlike some other languages (like Python), indentation in OrgLang is generally not semantically significant. It is used primarily for

readability and to reflect the structural hierarchy of nested blocks (e.g., inside `{}` or `[]`).

However, there are specific lexical rules where column position matters:

- The multiline comment marker `###` **must** start in the first column of the line.

Whitespace between tokens

Whitespace (spaces, tabs, and newlines) is used to separate tokens that would otherwise be joined. For example, `x : 42` requires whitespace around `:` if it were part of a larger word, but since `:` is a special symbol, it can often be used without spaces (e.g., `x:42`).

Certain symbols can be used without surrounding whitespace as they are recognized as distinct delimiters:

- `@`, `:`, `.`, `,`
- `(`, `)`, `[`, `]`, `{`, `}`

While not strictly required for these symbols, using whitespace is encouraged for visual clarity.

Identifiers and keywords

Identifiers

Identifiers (also referred to as names) are used to name variables, functions, and resources. In OrgLang, identifiers have a very flexible structure, allowing many symbols that are typically reserved for operators in other languages.

An identifier must start with a letter (case-sensitive `a-z` or `A-Z`), an underscore (`_`), or any of the following symbols:

- `!`, `$`, `%`, `&`, `*`, `-`, `+`, `=`, `^`, `~`, `?`, `/`, `<`, `>`, `|`

After the first character, an identifier can contain any combination of letters, underscores, digits (`0-9`), and the symbols listed above.

However, identifiers **cannot** contain the following structural delimiters:

- `@`, `:`, `.`, `,`, `;`, `(`, `)`, `[`, `]`, `{`, `}`

Examples of valid identifiers:

- `variable_name`
- `isValid?`
- `++count`
- `>>`
- `my-module`
- `$price`

Restricted Names: Identifiers that match any of the language's [Keywords](#) are reserved and cannot be used as variable names.

[NOTE] Since digits are allowed in identifiers but an identifier cannot *start* with a digit, the lexer can easily distinguish between numeric literals and names. For example, `42x` is not a valid identifier, but `x42` is.

Keywords

The following identifiers are reserved as keywords and have special meaning in the OrgLang language. They cannot be used as ordinary identifiers:

- `true` : Boolean truth value.
- `false` : Boolean falsehood value.
- `resource` : Used in resource definitions.
- `this` : Refers to the current function or block (used for recursion).
- `left` : Predefined name for the left operand in a binary operator.
- `right` : Predefined name for the right operand (or the operand of a prefix operator).

Reserved classes of identifiers

Literals

String literals

In OrgLang, strings are sequences of characters used for text representation. They can be defined as single-line or multiline literals.

Single-line Strings

Single-line strings are enclosed in double quotes (""). They must begin and end on the same line.

```
message : "Hello, OrgLang!";
```

Multiline Strings (DocStrings)

Multiline strings are enclosed in triple double quotes ("""). They can span multiple lines and are designed for large blocks of text or documentation.

To keep the source code clean, multiline strings automatically strip **common leading whitespace** (indentation) from all non-empty lines. The amount of whitespace removed is determined by the line with the least indentation.

```
# The resulting string will have no leading spaces
# on "Line 1" and "Line 2"
doc : """
    Line 1
    Line 2
""";
```

[NOTE] Leading and trailing blank lines (usually surrounding the delimiters) are also stripped.

Strings as Tables

A fundamental design choice in OrgLang is that **Strings are semantically Tables** (ordered lists of characters).

- **Indexing:** A string can be indexed by integers starting at `0`. Accessing `s.0` returns the first character.
- **Table Properties:** Because they are tables, operators like `->` for iteration treat strings as a stream of characters.

String Length and Orthogonality

Strings behave like tables when used in arithmetic contexts:

- **Numeric Value:** When used with arithmetic operators (like `+` or `-`), a string evaluates to its **length**.
- **No Concatenation:** Unlike many languages, the `+` operator does **not** concatenate strings. Instead, it adds their lengths (or a string's length to a number).

```
s1 : "ABC";
s2 : "DE";
res : s1 + s2; # res is 5 (3 + 2)
```

To join strings, use interpolation or specialized table operations (to be defined in the standard library).

Numeric literals

In OrgLang, all first-class numeric literals (integers and decimals) are designed to be implemented with **arbitrary precision**. This means that, by default, numbers are not limited by the standard 32-bit or 64-bit constraints of the underlying hardware, allowing for exact computations with very large or very precise values.

[NOTE] While the language semantics favor arbitrary precision, future versions of the compiler will introduce support for specific machine types (like `int`, `long`, `float`, `double`) as internal optimizations. These will be used when the compiler can prove that the range and precision requirements are satisfied by the more efficient machine representations.

Integer literals

Integer literals are represented as a sequence of one or more digits (`0-9`).

Signed Integer Literals

An integer literal can be preceded by an optional sign character (`+` or `-`).

[WARNING] To be treated as a single numeric literal, there **MUST NOT** be any whitespace between the sign and the digits.

- `-42` : A single integer literal token with value negative 42.
- `- 42` : Two tokens: the unary negation operator `-` followed by the integer literal `42`.

While the external behavior might often be similar, the distinction is important for the binding power of operators and lexer-level identification of values.

Decimal literals

Decimal literals represent non-integer numbers using a fixed decimal point notation. In OrgLang, these are distinct from the “floating point” types found in many other languages because they are designed for arbitrary precision and avoid the precision loss typical of binary floating point representations.

Decimal Syntax

A decimal literal consists of an integer part and a fractional part separated by a dot (`.`).

- **Digits on both sides:** For an unsigned number starting with a digit, there must be at least one digit on both sides of the dot (e.g., `3.14`).
- **Lexical Distinction:** A number followed by a dot without a subsequent digit (e.g., `1.`) is lexically interpreted as an [Integer literal](#) followed by the [Dot operator](#).
- **No Scientific Notation:** OrgLang does not currently support scientific notation (e.g., `1e10`) in its literal syntax.

Signed Decimal Literals

Like integers, decimal literals can be preceded by an optional sign (`+` or `-`) with no intervening whitespace.

```
pi : 3.14159;
negative_small : -0.0001;
positive : +1.0;
```

Rational literals

Rational literals represent exact fractional numbers using a syntax that clearly separates the numerator and denominator. This ensures precision in calculations involving fractions.

Rational Syntax

A rational literal is formed by two integer literals separated by a forward slash (`/`).

- **Numerator and Denominator:** Both the numerator and the denominator are integer literals (which can be positive or negative, as defined in [Integer literals](#)).
- **No Whitespace:** There must be no whitespace between the numerator, the slash, and the denominator.

- **Zero Denominator:** A zero in the denominator is syntactically valid but will result in a runtime error (division by zero) during evaluation.

Examples

```
one_half : 1/2;
three_quarters : 3/4;
negative_fraction : -1/2;
large_fraction : 123456789/987654321;
```

Boolean literals

Boolean literals represent truth values and correspond directly to the keywords `true` and `false`.

- `true`: Represents a true condition or success.
- `false`: Represents a false condition or failure.

These are the only two values of the Boolean type. They are used in conditional expressions (`?`, `?:`) and logical operations (`&&`, `||`, `!`).

Table literals

Tables (also referred to as Lists) are the primary data structure in OrgLang. There is a single, unified model for table construction, whether using blocks or operators.

Construction: Blocks and Commas

A table can be constructed in two ways that produce the same semantic object:

1. **Blocks** (`[]`): Square brackets group a sequence of statements or expressions into a Table. Elements are typically separated by whitespace.
2. **The Comma Operator** (`,`): The comma is a binary operator that creates or extends a Table.

Because `[]` evaluates its contents and collects the results into a new Table, using commas inside brackets results in **nesting**.

```

# A simple table
t1 : [1 2 3];
[...]
# The same table using commas (outside brackets)
t2 : 1, 2, 3;
[...]
# NESTED: [1, 2, 3] creates a Table containing a Table
t3 : [1, 2, 3]; # Result is [[1 2 3]]

```

Implicit Indexing vs. Bindings

A Table consists of a sequence of elements. These elements are categorized into two types:

- **Bindings:** Pairs created using the binding operator (`:`). These are accessed by their key and **do not** consume a numeric index slot.
- **Positional Elements:** All other expressions. These are assigned implicit numeric indexes (`0, 1, 2...`) based on their order of occurrence among other positional elements.

Mixed Content and Indexing

When a Table contains both bindings and positional elements, the numeric indexes skip the bindings.

```

mixed : [10, "status" : "active", 20];
[...]
val0 : mixed.0;          # 10
status : mixed."status"; # "active"
val1 : mixed.1;          # 20 (NOT mixed.2)

```

Tables as Blocks

Since every Org file is itself a Table, the rules for table literals apply to the top-level structure of a program. A file containing `a:1; b:2;` is a Table where `a` and `b` can be accessed by name.

Laziness in Tables

Values within a table are **lazy by default**. They are represented as thunks and are only evaluated when accessed (e.g., using the `.` or `?` operators).

```
computation : [1 + 1, 2 * 2]; # Expressions are not  
evaluated yet  
result : computation.0;      # 2 (evaluation happens here)
```

Operators

In OrgLang, almost every operation and structural construct is modeled as an **operator**. The language is designed to be highly orthogonal, with a minimal set of core rules that govern how these operators interact within expressions. Unlike many traditional languages that distinguish between operators, functions, and control structures, OrgLang treats nearly everything—from arithmetic to resource management and conditional evaluation—as an expression driven by operators.

Philosophy and Mechanics

OrgLang operators are strictly **unary** (prefix) or **binary** (infix). This strictness simplifies the language's grammar and execution model but introduces a different way of thinking about computation:

- **Binding Power:** The behavior of an expression is determined by the binding power (precedence) of its operators. Operators with higher binding power “pull” operands closer than those with lower power.
- **Everything is an Expression:** Operators don't just “perform actions”; they transform values and return new ones. This allows for deeply nested and highly expressive chains of computation.

Limitations and Patterns

The limitation to unary and binary forms (maximum of two operands) may seem restrictive compared to the variety of arities found in other languages. However, OrgLang overcomes this through several powerful patterns:

- **Tables as Parameters:** To pass multiple values to an operation that only accepts one operand (like a unary function call), those values are grouped into a [Table literal](#). The operation then extracts exactly what it needs from the table.

- **Currying:** Binary operators can be used to “partially apply” data. An expression like `a op b` can return a new thunk or function that is “ready” to take more data later.
- **Abstractions:** Simple operators can be composed and bound to names, creating high-level abstractions that behave like complex built-in features in other languages.

By embracing these patterns, OrgLang achieves a high degree of expressiveness while maintaining a structurally simple core.

Arithmetic operators

Arithmetic operators perform standard mathematical calculations. In OrgLang, these operators are designed to work with arbitrary-precision [Numeric literals](#).

OPERATOR	NAME	ARITY	DESCRIPTION
+	Addition	Binary	Returns the sum of two numbers.
-	Subtraction	Binary	Returns the difference between two numbers.
-	Negation	Unary	Returns the additive inverse of a number.
*	Multiplication	Binary	Returns the product of two numbers.
/	Division	Binary	Returns the quotient of two numbers.
**	Power	Binary	Returns the left operand raised to the power of the right operand (Right-associative).

[NOTE]

IMPLICIT COERCION

: Any arithmetic operator can be applied to [Table literals](#) and [Strings](#), in which case their `size` is used as the numeric value. Additionally, [Boolean literals](#) are coerced to numbers: `true` is treated as `1`, and `false` as `0`.

Bitwise operators

OrgLang supports standard bitwise operations for integers:

OPERATOR	NAME	ARITY	DESCRIPTION
&	Bitwise AND	Binary	Returns a number where each bit is <code>1</code> only if the corresponding bits of both operands are <code>1</code> .
\	Bitwise OR	Binary	Returns a number where each bit is <code>1</code> if at least one of the corresponding bits of the operands is <code>1</code> .
^	Bitwise XOR	Binary	Returns a number where each bit is <code>1</code> if the corresponding bits of the operands are different.
~	Bitwise NOT	Unary	Returns the bitwise complement (e.g., <code>~ 0 = -1</code>).
<<	Left Shift	Binary	Shifts the left operand's bits to the left by the number of positions specified by the right operand.
>>	Right Shift	Binary	Shifts the left operand's bits to the right by

OPERATOR	NAME	ARITY	DESCRIPTION
			the number of positions specified by the right operand.

Example:

```
10 & 2 // 2
10 | 5 // 15
10 ^ 5 // 15
~0      // -1
1 << 2 // 4
8 >> 1 // 4
```

Comparison operators

Comparison operators compare two values and always return a [Boolean literal](#) (`true` or `false`). OrgLang supports standard comparison operations, as well as automatic [type coercion](#) (e.g., comparing a string length to an integer).

[NOTE]

IMPLICIT COERCION

: Comparison operators follow the same coercion rules as [Arithmetic operators](#): Tables and Strings use their size, and Booleans are treated as `1` (`true`) or `0` (`false`).

OPERATOR	NAME	ARITY	DESCRIPTION
=	Equal to	Binary	Returns <code>true</code> if the operands are equal.
<> , ~=	Not equal to	Binary	Returns <code>true</code> if the operands are not equal.
<	Less than	Binary	Returns <code>true</code> if the left operand is strictly less than the right operand.
<=	Less than or equal to	Binary	Returns <code>true</code> if the left operand is less than or equal to the right operand.
>	Greater than	Binary	Returns <code>true</code> if the left operand is strictly greater than the right operand.
>=	Greater than or equal to	Binary	Returns <code>true</code> if the left operand is greater than or equal to the right operand.

[WARNING]

COMPARISON CHAINING

: Since every comparison returns a Boolean, the result of a chain (e.g., `x < y < z`) is the result of the **last comparison** in the chain. This differs from languages where such a chain might be shorthand for `(x < y) && (y < z)`.

Boolean operators

Boolean operators are used to perform logical calculations.

OPERATOR	NAME	ARITY	DESCRIPTION
!	Logical NOT	Unary	Returns the logical negation (e.g., <code>! 0 = 1</code>).
<code>~</code>	Bitwise NOT	Unary	Returns the bitwise complement (e.g., <code>~ 0 = -1</code>).
<code>&&</code>	AND	Binary	Short-circuit logical AND (returns <code>true</code> only if both are <code>true</code>).
<code>\ \ </code>	OR	Binary	Short-circuit logical OR (returns <code>true</code> if at least one is <code>true</code>).
<code>&</code>	Logical AND	Binary	Non-short-circuit logical AND.
<code>\ </code>	Logical OR	Binary	Non-short-circuit logical OR.
<code>^</code>	Logical XOR	Binary	Returns <code>true</code> if exactly one of the operands is <code>true</code> .

[NOTE]

TRUTHINESS

: Boolean operators can be applied to [Table literals](#) and [Strings](#).

They follow a “size-based” truthiness rule: a size of `0` is treated as `false`, and every other value (size `> 0`) is treated as `true`.

Conditional operators

Conditional operators allow for selection and branching within expressions without traditional `if/else` statements.

OPERATOR	NAME	ARITY	DESCRIPTION
.	Dot Access	Binary	Static/ Positional access to a Table's elements or keys.
?	Selection Access	Binary	Conditional or dynamic selection from a Table. Evaluation- driven.
??	Error Check	Binary	Returns the right operand if the left operand is an Error ; otherwise, returns the left operand.
?:	Elvis Operator	Binary	Returns the right operand if the left operand is "falsy" (false, Error, or an empty Table/ String); otherwise, returns the left operand.

Resource operators

Resource operators manage the lifecycle and data flow of [Resources](#).

Resource Instantiation (@)

The prefix `@` operator is used to instantiate a resource. When applied to a resource name or literal, it executes the resource's `setup` block and returns a **Resource Instance**.

```
# Instantiate stdout
@stdout
```

Data Flow (->)

The binary `->` operator drives data from a source (left operand) to a sink (right operand).

- **Source -> Sink:** Drives all data from the source into the sink until completion.
- **Iterator -> Function:** Creates a new projection (map) that will process elements lazily.

```
# Send a string to stdout
"Hello" -> @stdout;
[REDACTED]

# Send input through a transform to output
@stdin -> { args * 2 } -> @stdout;
```

Balanced Data Flow (-<)

The binary `-<` operator performs a balanced dispatch of data. It sends each element from the left source to exactly **one** of the available sinks in the Table on the right side, typically using a round-robin or load-balancing strategy.

- **Load Balancing:** If the right operand is a Table of sinks, elements are distributed among them.
- **Degeneration to ->:** If the right operand contains only one sink, it behaves identically to the basic data flow operator (`->`).

```
# Distribute tasks between two workers
@tasks -< [worker1 worker2];
```

Join Data Flow (-<>)

The binary `-<>` operator acts as a synchronizing barrier. It is used to merge multiple data streams into a single flow of coordinated packets.

- **Synchronization:** It waits for **every source** in the left Table to produce at least one element.

- **Aggregation:** Once one element is received from each source, it combines them into a single Table and sends that Table as a single “pulse” to the right operand.

```
# Synchronize data from two sensors before processing
[sensor1 sensor2] -<> processor;
```

Assignment operators

In OrgLang, assignment is strictly an operation that binds a value to a name within a [Table](#).

OPERATOR	NAME	DESCRIPTION
:	Binding	Binds the result of the right expression to the name specified on the left.

[NOTE]

EXTENDED ASSIGNMENT

: OrgLang reserves the following operators for extended assignment (modification of existing bindings). These are **not yet implemented** in the current runtime.

OPERATOR	DESCRIPTION	EXAMPLE
<code>:+</code>	Addition and Assignment	<code>x :+ 2</code>
<code>:-</code>	Subtraction and Assignment	<code>x :- 1</code>
<code>::*</code>	Multiplication and Assignment	<code>x ::* 3</code>
<code>:/</code>	Division and Assignment	<code>x :/ 4</code>
<code>:%</code>	Modulo and Assignment	<code>x :% 5</code>
<code>++</code>	Increment and Assignment	<code>++ x</code>
<code>--</code>	Decrement and Assignment	<code>-- x</code>
<code>:>></code>	Right Shift and Assignment	<code>x :>> 5</code>
<code>:<<</code>	Left Shift and Assignment	<code>x :<< 5</code>
<code>:&</code>	AND and Assignment	<code>x :& y</code>
<code>:^</code>	XOR and Assignment	<code>x :^ y</code>
<code>:\\ </code>	OR and Assignment	<code>x :\\ y</code>
<code>:~</code>	Bitwise NOT and Assignment	<code>x :~ 1</code>

Operator definitions

OrgLang allows for the definition of custom operators and the refinement of existing ones using the **Binding Power** syntax. This syntax defines the left and right binding powers, determining the operator's precedence and associativity.

```
# Define a unary operator with prefix power 100
! : 100{ ... };

# Define a binary operator with left power 50 and right
# power 60
op : 50{ ... }60;
```

When an operator is called, the expression within the braces is evaluated. The operands are made available via `left` and `right`.

- `left`: The left operand (for binary operators). For unary (prefix) operators, this is typically `Error` or `NULL`.
- `right`: The right operand (for binary operators) or the single operand (for unary operators).
- `this`: A reference to the operator function itself (useful for recursion).

[WARNING]

STRICT BINDING POWER SYNTAX

: When defining custom binding powers, there **must not be any whitespace** between the number and the brace.

- CORRECT

```
: op : 50{ ... }60;
```

- INCORRECT

```
: op : 50 { ... } 60;
```

Operators on operators

OrgLang provides higher-order operators that allow for the functional construction of logic by combining or specializing existing operators.

The `o` (compose) operator

The binary `o` operator performs **Functional Composition**. it merges two operators into a single, unified transformation.

- **Sequence:** In the expression `h : g o f`, the output of the right operator (`f`) becomes the input of the left operator (`g`).
- **Optimization:** The runtime attempts to fuse these operations into a single execution step to minimize intermediate overhead.

Arity-based Composition Rules (`h : g o f`):

The behavior of the composed operator `h` depends on the arity of `g` and `f`. The general rule is that the result of `f` always populates the `right` slot of `g`, and if `g` is binary, it retains the original `left` operand.

- **Unary `g` o Binary `f`:** `h` is a binary operator. `h(left, right)` evaluates as `g(f(left, right))`.
- **Binary `g` o Unary `f`:** `h` is a binary operator. `h(left, right)` evaluates as `g(left, f(right))`. This effectively uses `f` to pre-process the “main” argument while preserving the context in `left`.
- **Binary `g` o Binary `f`:** `h` is a binary operator. `h(left, right)` evaluates as `g(left, f(left, right))`.
- **Unary `g` o Unary `f`:** `h` is a unary operator. `h(right)` evaluates as `g(f(right))`.

```
# Compose increment and double
inc : { right + 1 };
double : { right * 2 };
inc_and_double : double o inc;

result : inc_and_double 5; # 12
```

The `|>` (partial application) operator

The binary `|>` operator, also known as the **Left Injector**, performs **Partial Application**. It “anchors” a value into the `left` slot of an operator, returning a new unary operator.

- **Specialization:** It allows you to create specialized versions of binary operators by fixing one of the operands.
- **Left Binding:** The value on the left of `|>` is bound to the `left` parameter of the operator on the right.

```
# Create a specialized 'add 10' function
add_ten : 10 |> +;

result : add_ten 5; # 15
```

Delimiters

Delimiters are structural symbols used for grouping expressions, constructing data structures, and defining blocks of code.

Parentheses ()

Parentheses are primarily used to **group expressions** and override the default precedence of operators.

```
res : (1 + 2) * 3; # 9
```

They are also used in function calls, although functionally `f(x)` is just `f` applied to the expression `(x)`.

Square Brackets []

Square brackets are used to construct **Table literals**. They group a sequence of expressions, evaluate them, and collect the results into a new Table.

```
list : [1 2 3];
nested : [[1 2 3]];
```

Braces { }

Braces are used to define **function bodies** and create **Operators**. The code inside braces is not executed immediately; instead, it is wrapped in an Operator (or thunk) that is evaluated when called.

```
# A simple function
add : { left + right };

# A thunk (parameter-less function)
thunk : { 1 + 1 };
```

Data model

The Data Model defines the fundamental entities and their relationships within OrgLang. It describes how information is represented, organized, and manipulated by the runtime. OrgLang is built on a foundation of extreme orthogonality and high-level abstractions, where complex behaviors emerge from the interaction of a small set of primitive types and universal operators.

Values and types

In OrgLang, information is represented as **Values**. A Value is a piece of data that can be bound to a name, passed as an argument to an operator, or returned as the result of an expression.

The language uses a **Dynamic Typing** model. This means that variables (bindings) do not have types; only the Values themselves carry type information. A variable can hold a Number at one point and a Table later in the execution.

Values vs. Objects

Unlike many object-oriented languages, OrgLang does not strictly distinguish between “primitive values” and “objects.” Every entity, from a simple Integer to a complex Resource Instance, is a first-class Value. Even Errors and Operators are treated as Values that can be manipulated and stored.

First-Class Expressions

Because OrgLang is built on a late-binding, lazy evaluation model, any piece of code enclosed in braces `{ }` is itself a Value—an **Operator**. This allows logic to be passed around as data, forming the basis for the language’s “Compositional” nature.

Extreme Orthogonality

A hallmark of OrgLang values is their predictable behavior across different operators. For instance, the addition operator `+` is defined for all types:

- Adding two Numbers produces their sum.
- Adding a Table to a Number uses the Table’s size.
- Adding two Tables returns the sum of their sizes.

This consistency reduces the need for “special cases” and allows for highly generic code.

The standard type hierarchy

OrgLang organizes its types into a logical hierarchy. While the runtime may implement these as a flat set of structures for performance, semantically they follow this inheritance pattern:

- Expression
 - Error
 - Name
 - Table
 - String
 - Number
 - Integer
 - Rational
 - Decimal
 - Boolean
 - Operator
 - Unary
 - Binary
 - Nullary
 - Resource

Special names

main

One of the special names is `main`. It is a special name because it is the entry point of the program. An org executable will look for a key named `main` and execute it. If `main` is not found, the program will exit with an error. The `main` key *must* be in the compiled org file, i.e., not in a submodule. It can be a function or an expression.

Execution model

The Execution Model describes how OrgLang programs are evaluated, how names are resolved, and how state is managed over time. The model is centered around the concept of **Persistent Tables** and **Lazy Evaluation**.

Naming and binding

In OrgLang, naming is not a separate storage mechanism but a structural property of [Tables](#).

Everything is a Table

Every scope in OrgLang—whether it's the global file scope, a code block `{ }`, or a module loaded from another file—is semantically a Table. When you perform an assignment using the binding operator `:`, you are performing a key-value insertion into the **Current Table**.

Dynamic Binding and Shadowing

Bindings are resolved dynamically based on lexical scope. When an identifier is evaluated, the runtime looks it up in the current table. If not found, it traverses upward through parent tables (e.g., from an operator's internal scope to the file's global scope). If you assign to a name that already exists in the current scope, the new value shadows (updates) the previous binding.

Evaluation of Bindings (Laziness)

A core feature of OrgLang is that table entries are **Lazy** by default. When you bind an expression to a name, the expression is wrapped in a “thunk” and stored. Evaluation only occurs when the name is explicitly accessed via the [Selection Access](#) operator.

```
x : 1 + 2; # 'x' stores the expression { 1 + 2 }
y : x;      # 'y' now also stores the same thunk
result : x; # Accessing 'x' triggers evaluation, result
            becomes 3
```

Closures and Scoping

In OrgLang, an **Operator** (lambda) is a first-class value that encapsulates logic and a reference to its birth environment. This environment is semantically a **Table**, allowing for the creation of closures.

The Scope-Table Relationship

While an operator's body is written as a sequence of expressions, the runtime manages the internal execution context as a **Table**.

- **Contextual Storage:** Every binding created within an operator (e.g., `x : 10;`) is stored in a local Table.
- **Expression Return:** Evaluating an operator executes all internal expressions in sequence, but the final value of the operator is only the result of the **last expression**.
- **Parent Linkage:** Every operator carries a pointer to the Table in which it was defined (the parent scope).

Lexical Closures

Closures are a natural consequence of OrgLang's **Lazy Evaluation** and **Parent Linkage**.

- **Capture by Reference:** When a nested operator is defined, it “captures” its parent's Table.
- **Thunk Preservation:** Because assignments are stored as thunks, a closure does not capture a snapshot of a value; it captures the **expression** bound to the name in the parent Table.
- **Arena Persistence:** The memory for captured Tables is managed within the **Arena**. As long as a closure is “reachable” (alive), its parent's Arena/Table remains intact.

```

# Closure Example: A power function generator
power_of : {
    exponent : right;      # Bound in the parent table
    { left ** exponent }; # This inner operator captures
    'exponent'
};

square : 2 |> power_of;    # 'square' is a closure where
                           exponent is 2
result : 4 -> square;     # 16

```

Resolution Logic

When an identifier is evaluated inside a closure:

1. Check the immediate local slots (`left`, `right`).
2. Check the current operator's local Table (bindings made inside the `{ }`).
3. Traverse to the **Parent Table** stored in the operator's environment.
4. Repeat until the **Global Table** (File Scope) is reached.

Errors

Errors in OrgLang are not “exceptions” that interrupt the flow of control; they are **First-Class Values** that participate in the data flow.

Error Propagation

Most operators in OrgLang are “Error-Aware.” If any operand of a binary or unary operation is an Error value, the operator does not perform its standard calculation. Instead, it immediately returns the Error value. This allows errors to propagate naturally through complex expressions until they reach a handler or the program’s output.

Error Generation

Crucially, `Error` is not a literal in OrgLang. You cannot write `x : Error` or `Error + 1` in your source code, as the word `Error` is a type name, not a value constructor. To force the generation of an Error value, you must perform an operation that is mathematically or logically invalid.

```
# Forcing an error through division by zero
val : 1 / 0; # 'val' now holds an Error value
```

Error Handling

Specifically designed operators like `??` (Error Check) and `?:` (Elvis) allow the programmer to detect and recover from Error values. These are the only operators that do not automatically propagate errors from their left operand.

```
# Propagating an error generated from invalid math
( (1/0) + 1 ) * 2; # Returns Error

# Handling an error
val : (1/0) ?? 0; # Returns 0
```

Terminal Signaling

If an Error value is returned by the `main` entry point or remains as the result of a top-level expression, the runtime typically signals this to the user via the system's standard error stream (stderr).

Arithmetic conversions

Arithmetic expressions in OrgLang are designed to be highly predictable and permissive, adhering to the principle of **extreme orthogonality**. Arithmetic operators (+, -, *, /, %) always aim to return a numeric value (Integer, Rational, or Decimal) by coercing their operands if necessary.

Coercion of Non-Numeric Types

When an arithmetic operator is applied to a non-numeric type, it is automatically coerced into a Number before the operation is performed:

- **Tables and Strings:** Coerced to their **size** (the number of elements or characters).
- **Booleans:** Coerced to 1 for true and 0 for false.

```
# Extreme orthogonality examples
"Hello" + 1;      # 5 + 1 = 6
[10 20 30] * 2;   # 3 * 2 = 6
true + true;       # 1 + 1 = 2
```

Division Rules

OrgLang handles division between Integers with special care to maintain precision without prematurely forcing floating-point representation.

- **Exact Division:** If one Integer divides another perfectly with no remainder, the result is an **Integer**.
- **Inexact Division:** If there is a remainder, the result is a **Rational**.

```
result1 : 4 / 2; # result1 is Integer 2
result2 : 3 / 2; # result2 is Rational 3/2
```

Numeric Promotion

When operations involve different numeric subtypes, the result is promoted to the most general type:

- Operations involving a **Decimal** typically produce a **Decimal**.
- Operations involving a **Rational** and an **Integer** typically produce a **Rational**.

Atoms

In the Lisp tradition, **Atoms** are the fundamental, indivisible building blocks of OrgLang expressions. An atom represents a specific value or a name that cannot be further broken down by the parser without changing its meaning.

Lexical Space and Separation

OrgLang uses whitespace (spaces, tabs, newlines) to separate tokens. The rules for where space is required, optional, or forbidden are critical to distinguishing between atoms and operators.

- **Mandatory Spaces:** Space is required to separate two atoms that would otherwise merge into a single identifier (e.g., `x y` is two names, `xy` is one).

- **Optional Spaces:** Spaces are optional around delimiters (`(()`, `[]`, `{ }`, `,`, `;`) and structural operators (`.`, `:`, `@`, `->`). For example, `(1+1)` is equivalent to `(1 + 1)`.

- **Forbidden Spaces:**

- **Signed Numbers:** There must be **no space** between a leading sign and the digits for it to be parsed as a negative or positive number atom (e.g., `-1` is a Number, `- 1` is the unary negation operator applied to `1`).
- **Binding Power:** There must be **no space** between an integer and the braces when defining custom binding powers (e.g., `700{ ... }701`).

Names (Identifiers)

Names are tokens used to refer to bindings in a Table.

- **Characters:** Names can contain letters, digits, underscores, and most symbols (e.g., `isValid?`, `counter_1`, `set!`).
- **Start Rule:** A name cannot start with a digit (which would initiate a Number).

- **Exclusions:** Symbols used as delimiters or structural operators ((), [], { }, , ., :, @, ;) cannot be part of a name.

Literals Representation

Literals are atoms that represent fixed values.

Numbers

Numbers are represented in three subtypes:

- **Integers:** Sequences of digits, optionally preceded by a sign (42 , -10).
- **Decimals:** Digits containing a decimal point (3.14 , -.5).
- **Rationals:** Represented as a ratio of two integers (2/3).

Booleans

The keywords `true` and `false` are the only two Boolean atoms.

Strings

Strings are represented by text enclosed in double quotes:

- **Simple:** "Single line string".
- **Multiline:** """Triple quotes for multiline text""".

Strings are semantically **Tables** where each character is a value indexed by its position.

Atoms in Tables

One of the most important aspects of Atoms in OrgLang is how they interact with Table literals []. While the source code of a program is itself a Table where expressions are separated by semicolons ; , a Table literal uses **space** as its primary separator.

Because spaces are also used to separate atoms within an expression (like `1 + 1`), the use of space inside a Table literal can

be ambiguous. In OrgLang, the **space in a Table literal acts as an element separator**, effectively terminating the current expression.

- **Atomic Gathering:** Within `[]`, the runtime treats each space-separated sequence as a distinct element if not explicitly grouped.
- **Binding Greediness:** The binding operator `:` inside a Table literal is “greedy” and attaches the name on its left to the immediate next atom on its right. It does not automatically consume subsequent atoms if they are separated by spaces.

```
# Likely an error if you wanted 'a' to be 2:  
# 1. A binding/pair (a: 1)  
# 2. An operator (+)  
# 3. An integer (1)  
# Resulting Table: [0: +, 1: 1, a: 1]  
table_split : [a: 1 + 1];  
  
# Correct: results in a Table with one specific entry  
# The parentheses group '1 + 1' into a single Atomic  
# Expression  
table_ok : [a: (1 + 1)];
```

Parenthesized forms

Parentheses `()` are used to group expressions to override precedence. A parenthesized expression is treated as a single atomic unit (an **Atomic Expression**) during the evaluation of the outer expression or when used as an element in a Table.

Advantages and Limitations

The atomic model of OrgLang provides a unique balance of simplicity and expressive power:

- **Advantage: Extreme Orthogonality:** Because every atom (even Operators and Errors) is a first-class Value, the language’s core operators work consistently across all data types.
- **Advantage: Structural Purity:** Every source file is semantically a Table literal, making the relationship between code and data perfectly transparent.

- **Limitation: Lexical Sensitivity:** The reliance on space as an element separator in Tables means developers must be mindful of grouping when mixing spaces and operators.
- **Limitation: Precedence Quirks:** Some operators, like unary negation, have lower precedence than exponentiation ($-1^{**}2 = -1$), which preserves mathematical convention but may surprise users coming from languages where unary operators are always highest.

Unary arithmetic and bitwise operations

Unary operators in OrgLang are prefix operators that associate with the immediate expression to their right. They follow the principle of extreme orthogonality, coercing non-numeric types to numbers when necessary.

Negation (-)

The unary negation operator reverses the sign of a numeric value.

- **Integers:** Returns a negative or positive Integer.
- **Decimals:** Returns a Decimal with the sign reversed.
- **Precedence Note:** Unary negation has **lower precedence** than exponentiation. This means `-1**2` is evaluated as `- (1**2)`, resulting in `-1`, which aligns with standard mathematical notation.

Increment and Decrement (++, --)

These operators perform primitive arithmetic addition or subtraction of 1.

- `++x` is semantically equivalent to `x + 1`.
- `--x` is semantically equivalent to `x - 1`.
- Note: These are prefix operators and do not have “postfix” variants in the core language.

Bitwise NOT (~)

The bitwise NOT operator returns the bitwise complement of a number.

- **Coercion:** Non-integers are coerced to Integers before the bitwise inversion occurs.
- **Result:** Always returns an Integer. For example, `~0` results in `-1` (using two’s complement representation).

Logical NOT (!)

The logical NOT operator performs a truthiness check and returns a Boolean.

- **Truthiness:** `0`, `none`, and empty tables/strings are typically considered falsey. All other values are truthy.
- **Result:** Returns `true` if the operand is falsey, and `false` otherwise.

Binary arithmetic operations

Binary arithmetic operations in OrgLang are designed to be permissive and mathematically intuitive. They operate on three numeric types: **Integer**, **Rational**, and **Decimal**.

Standard Operators

- **Addition (+)**: Performs numeric addition. If both operands are integers and the result is within integer range, results in an **Integer**. If any operand is a Decimal, results in a **Decimal**.
- **Subtraction (-)**: Performs numeric subtraction. Same promotion rules as addition.
- **Multiplication (*)**: Performs numeric multiplication.
- **Division (/)**: Follows specialized precision rules:
 - **Exact**: $4 / 2$ results in **Integer 2**.
 - **Inexact**: $3 / 2$ results in **Rational 3/2**.
 - **Decimal**: If either operand is a Decimal, the result is a **Decimal**.
- **Modulo (%)**: Returns the remainder of division. Typically used with Integers.

Exponentiation (**)

The exponentiation operator raises the left operand to the power of the right operand.

- **Precedence**: Higher than unary negation. Thus, -1^{**2} is parsed as $-(1^{**2})$ resulting in -1 .
- **Promotion**: Often results in a **Decimal** if the power is fractional or negative, unless the result can be exactly represented as an Integer or Rational.

Numeric Coercion

Following the principle of **extreme orthogonality**, binary operators automatically coerce non-numeric types into Numbers:

- **Tables/Strings:** Their `size` is used as the numeric value.
- **Booleans:** `true` becomes `1`, `false` becomes `0`.
- **Errors:** Propagate through the operation (the result of any arithmetic with an Error is an Error).

```
# Examples of extreme orthogonality
[10 20] + [30]; # 2 + 1 = 3
"abc" * 2;      # 3 * 2 = 6
true + 1;       # 1 + 1 = 2
```

Shifting operations

Shifting operations in OrgLang are bitwise operations that operate on numeric values, treating them as bit patterns.

Left Shift (`<<`)

The left shift operator moves the bits of the left operand to the left by the number of positions specified by the right operand.

- **Coercion:** Non-integers are coerced to Integers via their numeric representation or size.
- **Arithmetic Effect:** Shifting an integer left by N bits is equivalent to multiplying it by 2^N .
- **Result:** Always an Integer.

Right Shift (`>>`)

The right shift operator moves the bits of the left operand to the right by the number of positions specified by the right operand.

- **Coercion:** Standard numeric coercion applies.
- **Arithmetic Effect:** Shifting an integer right by N bits is equivalent to integer division by 2^N .
- **Result:** Always an Integer.

```
# Examples of shifting
1 << 3;      # 1 * 8 = 8
16 >> 2;     # 16 / 4 = 4
"abcd" << 1; # 4 << 1 = 8 (uses string size)
```

Binary bitwise operations

Binary bitwise operators in OrgLang perform bit-by-bit operations on their operands. Like other arithmetic-adjacent operators, they follow the principle of **extreme orthogonality**, coercing non-numeric types to Integers.

Bitwise AND (&)

Returns a number where each bit is `1` only if the corresponding bits of both operands are `1`.

- **Coercion:** Standard numeric/size coercion applies.

- **Result:** Always an Integer.

Bitwise OR (|)

Returns a number where each bit is `1` if at least one of the corresponding bits of the operands is `1`.

- **Result:** Always an Integer.

Bitwise XOR (^)

Returns a number where each bit is `1` if the corresponding bits of the operands are different.

- **Result:** Always an Integer.

```
# Examples of bitwise operations
5 & 3;      # 101 & 011 = 001 (1)
5 | 3;      # 101 | 011 = 111 (7)
5 ^ 3;      # 101 ^ 011 = 110 (6)
"abc" & 7; # 3 & 7 = 3 (uses string size)
```

Comparisons

Comparison operators in OrgLang are used to determine the relationship between two values. They primarily operate on numeric values but, following the **extreme orthogonality** principle, will coerce non-numeric types to their numeric equivalent (usually their size).

Comparison Operators

OPERATOR	DESCRIPTION	EXAMPLE
<code>=</code>	Equal to	<code>x = y</code>
<code><> , ~=</code>	Not equal to	<code>x <> y</code>
<code><</code>	Less than	<code>x < y</code>
<code><=</code>	Less than or equal to	<code>x <= y</code>
<code>></code>	Greater than	<code>x > y</code>
<code>>=</code>	Greater than or equal to	<code>x >= y</code>

- **Result:** All comparison operators return a **Boolean** (`true` or `false`).

- **Equality Limitation:** OrgLang does **not** allow the specialization of equality. There are no plans to support custom equality methods for user-defined tables or operators; equality behavior is baked into the runtime for core types.

- **Numeric Comparisons:**

- **Integer vs. Rational:** Comparison is **exact**. The Integer is treated as a Rational with a denominator of 1.
- **Decimal vs. Others:** If any operand is a **Decimal**, the other operand is promoted to a Decimal (inexact) before the comparison is performed.

- **Decimal vs. Decimal:** Compared based on their approximate numeric values.
- **Structural Comparisons:** For strings or tables, comparisons currently default to evaluating their **sizes**.

Comparison Chaining

OrgLang allows comparison operators to be chained together, such as `x < y < z`.

- **Semantics:** In a chain of comparisons, each operation is evaluated in sequence. However, unlike languages like Python (where `x < y < z` is `(x < y) and (y < z)`), OrgLang's current execution model evaluates the chain left-to-right, and the **result of the entire chain is the result of the last comparison**.
- **Example:** `3 < 5 < 2` would evaluate `3 < 5` (returning `true`, which is `1`), then evaluate `1 < 2`, resulting in `true`.

[WARNING] Because of how chaining works, users should be cautious. If a mathematical range check is desired, explicit logical ands should be used once available (e.g., `(x < y) && (y < z)`).

Numeric Coercion in Comparisons

As with arithmetic, non-numeric types are coerced to numbers before comparison:

- **Strings/Tables:** Coerced to their **length/size**.
- **Booleans:** `true` is `1`, `false` is `0`.

```
# Examples of orthogonal comparisons
"apple" > "pear"; # 5 > 4 = true
[1 2] = 2;          # 2 = 2 = true
true < 2;           # 1 < 2 = true
```

Boolean operations

Boolean operations in OrgLang are used to combine or invert boolean values. They follow the principle of **extreme orthogonality**, meaning they can be applied to any type by first determining its **truthiness**.

Truthiness Rules

In OrgLang, nearly every value can be evaluated in a boolean context.

- **Truthy:** Any non-zero Number, any non-empty Table (including non-empty Strings), and any active Resource or Operator.
- **Falsey:** The number `0`, empty Tables `[]` and empty Strings `""`.

Logical Operators

OPERATOR	DESCRIPTION	RESULT TYPE
<code>&&</code>	Logical AND	Boolean
<code>\ \ </code>	Logical OR	Boolean
<code>!</code>	Logical NOT	Boolean

- **Short-circuiting:** Per the language design, `&&` and `\|\|` are **short-circuiting** operators. If the result can be determined by the left operand (e.g., `false && ...` or `true \|\| ...`), the right operand is not evaluated.
 - *Implementation Note:* In the current prototype, these may behave strictly; however, programs should be written assuming short-circuiting behavior.
- **Negation (`!`):** Returns `true` if the operand is falsey, and `false` if it is truthy.

```
# Examples of boolean logic
(1 = 1) && (2 = 2); # true
0 || "hello";         # true (0 is falsey, non-empty string
                      is truthy)
! [ ];                # true (empty table is falsey)
! "non-empty";        # false
```

Access and Conditional Expressions

OrgLang provides powerful operators for accessing data and handling conditions, aligning with its table-centric and expression-based philosophy.

Dot Access (.)

The `.` operator provides static or positional access to a Table's elements.

- **Behavior:** It retrieves the value associated with a key from the Table on the left.
- **Keys:** The key on the right can be an identifier (string key) or an integer index.
- **Thunk Evaluation:** Accessing a value via `.` forces the evaluation of the thunk stored at that position (if it hasn't been evaluated yet).

```
data : [10, "x": 20];
val1 : data.0;    # 10
val2 : data.x;   # 20
```

Selection Access

The `?` operator acts as a conditional or dynamic lookup. It reverses the standard access order: `condition ? table`.

- **Behavior:** It evaluates the `condition` (left operand) and uses the result as a `key` to look up a value in the `table` (right operand).
- **Usage:** Commonly used with a table containing `true` and `false` keys to mimic an if-else expression.
- **Lazy Evaluation:** Since table values are lazy (thunks), only the selected branch is evaluated.

```
result : (x > 0) ? [true: "Positive" false: "Non-
positive"];
```

Elvis Operator (?:)

The `?:` operator is a “falsy” coalescing operator.

- **Behavior:** Returns the left operand if it is **true**. If the left operand is **falsey** (0, empty string or empty table), it returns the right operand.
- **Usage:** Providing defaults for potentially empty or invalid values.

```
name : input_name ?: "Guest"; # Use "Guest" if input_name  
is empty
```

Error Coalescing (??)

The `??` operator is a specialized “Error” coalescing operator, similar to Null Coalescing in other languages.

- **Behavior:** Returns the left operand if it is **not** an Error. If the left operand is an **Error**, it returns the right operand.

```
# If calc_value returns 0, result is 0. If it returns  
Error, result is 10.  
result : calc_value() ?? 10;
```

Lambdas (Anonymous Operators)

In OrgLang, functions are first-class values called **Operators**. Anonymous operators (lambdas) are defined using curly braces `{ ... }`.

Implicit Parameters

Every operator has two implicit parameters available within its scope:

- **right**: The primary argument. In a binary expression `(a op b)`, this is `b`. In a unary expression `(op x)`, this is `x`.
- **left**: The secondary argument. In a binary expression `(a op b)`, this is `a`. In a unary expression, `left` is bound to **Error**.

Defining Operators

```
# Unary operator (function)
square : { right * right };
result : square 4; # 16

# Binary operator
add : { left + right };
result : 4 add 5; # 9
```

Recursion (`this`)

The keyword `this` refers to the current operator itself, allowing for anonymous recursion.

```
factorial : {
  (right <= 1) ? 1;
  (right > 1) ? (right * this(right - 1));
};
```

[NOTE]

VARIABLE CAPTURE

: Currently, operators in OrgLang do not capture their lexical environment (closures). They are pure functions of their inputs (`left`, `right`) and global values.

Table Construction

Tables are the fundamental data structure in OrgLang. There are three primary ways to construct them, each serving a different syntactic purpose but resulting in the same underlying data type.

1. Block Constructor ([])

The square brackets create a new Table scope. Expressions inside are evaluated, and their results are collected into the Table.

- **Separator:** Elements can be separated by spaces or newlines.
- **Scope:** Variables defined inside [. . .] are local to that table.
- **Usage:** Defining data structures, lists, or scoped blocks of code.

```
data : [ 1 2 3 ];
config : [
    host: "localhost"
    port: 8080
];
```

2. Comma Operator (,)

The comma is a binary operator that constructs a Table from its operands. It behaves differently depending on whether the left operand is already a Table.

- **Atom, Atom:** Creates a new Table with two elements.
 - 1, 2 → [1 2]
- **Table, Atom:** Appends the right operand to the left Table.
 - [1 2], 3 → [1 2 3]
- **Atom, Table:** Creates a new Table where the right operand is the second element (nesting).
 - 1, [2 3] → [1 [2 3]]

- **Table, Table:** Appends the right Table as a *single element* to the left Table (nesting).

◦ [1 2], [3 4] → [1 2 [3 4]]

[!NOTE]

LEFT-ASSOCIATIVE

: Because the comma is left-associative, 1, 2, 3 is parsed as (1, 2), 3.

1. (1, 2) becomes [1 2]
2. [1 2], 3 becomes [1 2 3] This makes it efficient for constructing lists.

3. Source File (Implicit Table)

In OrgLang, **every source file is implicitly a Table**.

- **Behavior:** The top-level scope of a file behaves exactly like the inside of a [. . .] block.
- **Rationale:** This allows OrgLang programs to look like standard script files (sequences of statements) without needing to be wrapped in a main function or a global object.
- **Result:** Importing a file involves executing it and returning the resulting Table.

```
# my_module.org
x : 10;
y : 20;
# This file evaluates to a Table: [ x: 10, y: 20 ]
```

Comparison

METHOD	SYNTAX	PRIMARY USE CASE	SCOPE
Block	[...]	Data structures, local scopes	New local scope
Comma	a , b	Tuples, arguments	Current scope
File	(Implicit)	Modules, scripts	File-level scope

Evaluation Order

OrgLang is a **Lazy** language by default for its tables and assignments, but **Eager** (mostly) for function calls, with specific evaluation rules.

Standard Rule: Left-to-Right

In general, expressions are evaluated from left to right. This applies to function arguments, list elements, and most binary operators.

Exceptions

1. Lazy Assignment / Thunks The most critical exception in OrgLang is the **assignment operator** (`:`).

- **Rule:** The expression on the right side of an assignment is **NOT evaluated** at the moment of assignment.
- **Thunks:** Instead, it is stored as a “Thunk” (a suspended computation).
- **Evaluation:** The value is only computed when it is **used** (e.g., accessed via `.`, used in arithmetic, or printed).
- **Consequence:** The “type” of every assignment is effectively an **Expression** (or Thunk) until it is forced.

```
x : 1 / 0;    # No error here! The division is not
                performed yet.
y : x + 1;    # Still no error! 'y' is now a thunk
                depending on 'x'.
y -> @stdout; # Error happens HERE, when @stdout
                pulls 'y', which pulls 'x'.
```

2. Short-Circuiting Operators

- `&&` and `||`: The left operand is evaluated first. The right operand is evaluated **only if necessary**.
- `? , ?: , ??`: The condition (left) is evaluated first. Only the selected branch (right or fallback) is evaluated.

3. Right-Associative Operators For operators that are right-associative, the parsing groups them from right to left, but the

evaluation of operands typically still happens left-to-right before the operator function is called, unless they are lazy constructions.

- **Assignment (:**): `a : b : c` parses as `a : (b : c)`.

- **Exponentiation (**)**: `2 ** 3 ** 4` parses as `2 ** (3 ** 4)`.

4. Unary Operators Prefix operators (`-`, `!`, `~`) evaluate their single operand (to the right) before applying the operation.

Operator Precedence

OrgLang uses a **Pratt Parser** to handle operator precedence and associativity. The following table lists operators from highest binding power (tightest binding) to lowest.

PRECEDENCE	OPERATOR	DESCRIPTION	ASSOCIATIVITY
900	@, ~, !, -	Unary Prefix	N/A
800	. , ? , ?? , ?:	Access / Call / Conditional	Left
500	**	Exponentiation	Right
400	o , \ >	Composition / Injection	Left
300	* , / , & , %	Product / Bitwise AND	Left
200	+ , - , \ , ^ , << , >>	Sum / Bitwise OR/XOR/Shift	Left
150	= , <> , < , > , <= , >=	Comparisons	Left
140	&&	Logical AND	Left
130	\ \	Logical OR	Left
100	(User Defined)	Custom Operators	Left (Default)
80	:	Binding (Assignment)	Right
60	,	Comma (Table construction)	Left
50	-> , -< , -<>	Flow / Dispatch / Join	Left
0	;	Statement Terminator	N/A

[NOTE]

USER-DEFINED OPERATORS

: By default, custom operators fall into the **100** slot. This places them below standard arithmetic but above flow and assignment, minimizing ambiguity in pipelines.

Assignment

Assignment in OrgLang is fundamentally about **constructing Tables**. Since every scope is a Table, “assigning” a value to a name is equivalent to creating a **Pair** (Key-Value binding) within that Table.

Standard Assignment (:)

The colon operator `:` binds a value (the right operand) to a name (the left operand).

- **Semantics:** Creates a Pair `[Name Value]`. When occurring inside a Table constructor (including a source file), this Pair becomes a named entry in that Table.
- **Laziness:** As noted in [Evaluation Order](#), the value is stored as a **Thunk** and is not evaluated until used.

```
x : 10;      # Binds 'x' to a thunk returning 10
config : [
    port : 8080  # Binds 'port' to 8080 within the
                  'config' table
];
```

Extended Assignment (Reserved)

OrgLang reserves a set of operators for **extended assignment**, which combines an operation with assignment (modification).

[WARNING] These operators are
NOT YET IMPLEMENTED
in the current runtime. They are reserved for future versions to support in-place mutation or syntactic sugar.

OPERATOR	DESCRIPTION	EQUIVALENT TO (CONCEPTUALLY)
<code>;+</code>	Add and Assign	<code>x : x + y</code>
<code>; -</code>	Subtract and Assign	<code>x : x - y</code>
<code>; ~</code>	Concatenate and Assign	<code>x : x ~ y</code>
<code>; >></code>	Right Shift and Assign	<code>x : x >> y</code>
<code>; <<</code>	Left Shift and Assign	<code>x : x << y</code>
<code>++</code>	Increment (Prefix)	<code>x : x + 1</code>
<code>--</code>	Decrement (Prefix)	<code>x : x - 1</code>

[NOTE] Currently, `++` and `--` are implemented as **Arithmetic Operators** that return `value + 1` or `value - 1` without modifying the variable. Full mutation support is planned for a future release.

Resources Execution

Resources in OrgLang are the bridge between the pure, immutable world of your program and the mutable, effectful world outside (file system, network, screen). They are based on **Algebraic Effects**, meaning that “doing” something is separated from “interpreting” it.

The definition of a Resource is done with the `@:` operator. The left side must be a name, the right side a Table with the following keys:

- `next` : The step operator, a unary operator that is called when the resource is used.
- `create` : The init operator, a unary operator that is called when the resource is created.
- `destroy` : The close operator, a nullary operator that is called when the resource is closed.

The `create` operator may be nullary, if default arguments are provided.

[NOTE] The `next` operator is the only one that is not optional. If it is not provided, the resource will not be able to be used. The `create` operator is optional. If it is not provided, the resource will not be able to be used. The `destroy` operator is optional. If it is not provided, the resource will not be cleaned up after use. The `@:` operator is the only way to define a resource. This is just syntactic sugar.

Concept: Effect Reification

When you write to a file or print to the screen, you aren’t just calling a function. You are creating a **Resource Instance** that represents that interaction. The runtime then “interprets” this instance to perform the actual side effect.

Resource Primitives vs. Standard Library

It is important to distinguish between **Runtime Primitives** (implemented in C) and **Standard Library Resources** (implemented in OrgLang).

1. Runtime Primitive (@sys) Currently, OrgLang exposes a single, low-level primitive for system interaction: `@sys`.

- **Behavior:** It accepts a **Command List** where the first element is a string (e.g., “write”, “read”) and subsequent elements are arguments.
- **Purpose:** It acts as a direct bridge to C-level system calls.

[NOTE] The `@sys` primitive is the only way to interact with the system. It is not a resource, but a primitive that can be used to create resources. It will probably be removed in the future, and replaced by a set of more specialized primitives.

2. Standard Library Resources These are high-level abstractions built *using* the `@sys` primitive.

- `@stdout`: Wrapper around `["write" 1 ...] @ sys`.
- `@stdin`: Wrapper around `["read" 0 ...] @ sys`.

```
# How stdout is implemented (conceptually)
stdout @: [
    next: { ["write" 1 right -1] @ sys }
];
```

Planned Primitives: Future versions may introduce specialized primitives like `@file` or `@net` for better performance and type safety, reducing reliance on the generic `@sys`.

The `@:` Operator

You can define your own resources using the `@:` operator. A resource definition is a Table that acts as a blueprint, specifying how to handle the lifecycle of the effect.

Lifecycle Hooks:

- `create`: (Optional) Called when the resource is instantiated. Returns the initial state.

- `next`: (Optional) Called when data is pulled *from* the resource or pushed *into* the resource.
- `destroy`: (Optional) Called when the resource is closed or goes out of scope.

```
# A simple logger resource
Logger @: [
    next: {
        # 'right' is the data being pushed
        ["write" 1 ("LOG: $0\n" $ [right])]
        @ sys # Syscall to write to FD 1
    }
];
```

Instantiation (@)

To use a resource, you must **instantiate** it using the prefix `@` operator. This creates a live instance with its own state.

```
main: {"Hello" -> @Logger} # Push data to it
> H
> e
> l
> l
> o
```

In this simple example, the `Logger` resource is instantiated and the string `"Hello"` is pushed to it. The `next` operator is called with `right` being `"Hello"`. The `next` operator then calls the `write` syscall with `1` as the file descriptor and `"LOG: Hello\n"` as the data. The Logger resource is then destroyed since the source (the "Hello" string) is consumed.

The strange output is due to Strings being Tables, so every character is sent to the resource one by one. To achieve what is most likely intended, we should send the string inside a Table.

```
main: {[ "Hello" ] -> @Logger}
```

[NOTE] The `@` operator *must* be used inside operator definitions, i.e. `{...}`.

Flux Operators

Resources are primarily used with Flux operators to move data.

- -> (**Push**): Sends data to a resource's `next` function.
`source -> @stdout;`
- -< (**Dispatch**): load-balances data across multiple resources.
- -<> (**Join**): Synchronizes multiple resource streams.

Scoped Resources (@arena)

Scoped resources are a special pattern where a resource manages a memory arena for the duration of a flow. This allows for safe, high-performance heap allocation with deterministic teardown.

- **Arena:** The `@arena` resource creates a memory context. Any resources (like `@file` or `@net`) created “downstream” of an arena are tracked by it.
- **Middleware Pattern:** Arenas are often used as middleware in a flow chain to wrap the execution context.
- **Teardown:** When the flow ends or the scope is exited, the Arena automatically tears down all tracked resources in reverse order of creation.

```
# Use Arena to manage memory for a file operation
"data.txt" -> @file_reader -> @arena -> @process ->
    @stdout;
```

In this example, `@file_reader` allocates its buffer within `@arena`. When the flow completes, `@arena` is freed, automatically closing the file and reclaiming memory.

Memory management

In **OrgLang**, the Arena model is designed to be high-performance yet flexible enough to handle the non-linear growth typical of recursive functions and unpredictable data streams. By moving away from a single fixed block to a **Chained Page Strategy**, we ensure that the language remains “systems-ready” while keeping the “pointer-bumping” speed.

Here is how the architecture handles your specific concerns:

Handling Recursive Functions

Recursive functions present a unique challenge for Arenas because each call creates a new scope but depends on the parent’s data.

- **Sub-Arena Nesting:** Every time a `{ }` block is entered, the runtime doesn’t necessarily create a new heap; it simply marks a “checkpoint” in the current Arena.
- **Tail-Call Optimization (TCO):** Since OrgLang uses **Lazy Evaluation** and **Thunks**, the compiler can often detect if a recursive call is the final operation. In these cases, it performs a “jump” instead of a “push,” reusing the current Arena frame and preventing stack/memory overflow.
- **Closure Persistence:** If a recursive function returns a closure, that closure “holds onto” its birth Arena. The Arena remains alive until the closure itself is no longer reachable, ensuring captured variables (like a recursive state) don’t vanish prematurely.

Dealing with Unknown Memory Demands

Since the compiler cannot always predict how much memory a flow (`->`) will need, OrgLang employs a **Chained Page Allocation** strategy.

The Allocation Strategy: “The Multi-Page Arena”

Instead of a single `mmap` call for a giant block, the runtime works in **Pages** (typically 4KB or 2MB).

1. **Initial Page:** When a flow starts, the runtime requests one memory page from the OS via `@sys("mmap")`.
2. **Pointer Bumping:** Allocation happens by incrementing the `top` pointer.
3. **Page Overflow:** If `top + request > page_end`, the runtime allocates a **new page** and links it to the previous one (creating a linked list of pages).
4. **Bulk Deallocation:** When the flow ends, the runtime doesn’t free items individually; it iterates through the page list and releases every page back to the OS in one go.

Allocation Strategy Details

FEATURE	IMPLEMENTATION	BENEFIT
Growth	By-Demand Paging	Doesn't waste RAM upfront for small scripts; scales to GBs for large flows.
Small Objects	Internal Fragmentation	Since we don't free individual items, we don't care about "holes." We just keep moving forward.
Large Objects	Dedicated Pages	If you request <code>@mem(10MB)</code> , the runtime bypasses standard pages and maps a specific 10MB region for that object.
Cleanup	Reverse Teardown	Before releasing pages, the runtime walks the "Teardown Registry" to close file descriptors or sockets registered in that Arena.

Final Specification Note: The “Frame” Reset

For long-running streams (like `@stdin -> @stdout`), the scheduler can perform **Frame Resets**. Between pulses, if the runtime detects that a pulse's data has been fully consumed by all sinks, it can reset the Arena pointers for that specific branch, effectively providing “infinite” execution in a small, fixed memory footprint.

Operator Definitions

In OrgLang, operators are fundamentally **functions bound to symbolic names** within the current scope. This aligns with the language's philosophy that “Scope is a Table”.

Defining Custom Operators

To define a new operator (or override an existing one), you assign a **Function** to a string key corresponding to the operator symbol.

```
# Define a custom operator |+|
"|+" : {
    left + right + 1
};

# Usage
x : 10 |+| 20; # x = 31
```

- **Symbol:** The key must be a valid operator symbol string.
- **Arity:** The function receives `left` and `right` arguments for infix operators. For prefix operators, `left` is typically `null` or `error`.
- **Precedence:** Custom operators currently default to a precedence level of `100`, placing them below standard arithmetic but above assignment.

Execution Model

When the runtime encounters an operator expression `A op B`:

1. **Resolution:** It looks up the operator string `op` in the current **Scope** (Table).
2. **Dispatch:**
 - If a user-defined function is found, it is called with `A` and `B`.
 - If not found, it falls back to the **Standard Library** implementation (e.g., primitive `+`).

3. **Binding:** This dynamic dispatch allows for powerful DSL creation and localized operator behavior.

Module Definitions

Modules in OrgLang are simply source files `.org` that are evaluated as **Tables**.

Importing a Module

To load a module, you use the `@` operator with the `org` resource.

```
lib : "path/to/lib.org" @ org;
```

Module Execution Model

When a module is imported:

1. **Scope:** It is executed in its own **local scope**, isolated from the importer.
2. **Structure:** The entire file represents a data structure (Table), where entries are defined using the `:` operator.
3. **Result:** The module returns this **Table**, allowing access to all defined symbolic names.

lib.org:

```
helper : { left + 1 };
add_one : helper;
constant : 42;
```

main.org:

```
lib : "lib.org" @ org;
x : 10 -> lib.add_one; # x = 11
y : lib.constant; # y = 42
```

Full Grammar Specification

The following EBNF describes the valid sequences of tokens in OrgLang. This grammar does not capture operator precedence, which is handled dynamically by the runtime (Pratt Parser). It also does not enforce the “no space” rule for Binding Powers, which must be handled by the lexer/parser.

```

/* Lexical Tokens & Literals */
/* Note: Signed numbers must not have a space between the
sign and the number */
INTEGER      ::= ("-"?)? [0-9] +
DECIMAL      ::= ("-"?)? [0-9] + "." [0-9] +
RATIONAL     ::= INTEGER "/" INTEGER /* Syntactic sugar for
division expression */
STRING        ::= '"' [^"]* "'"
                  | '"""' .*? '"""'
BOOLEAN       ::= "true" | "false"

IDENTIFIER   ::= [a-zA-Z_!$%&*+\-=\^~?/<|>][a-zA-Z0-9_!$%&*+
\-=\^~?/<|>\.] *

/* Program Structure */
Program       ::= Statement*
Statement     ::= Expression (";")?

/* Expressions (Flat Structure) */
/* An expression is a sequence of operands and operators.
Precedence is dynamic. */
Expression    ::= Operand (Operator Operand)*

/* Operands */
Operand      ::= Literal
                  | Identifier
                  | Keyword
                  | Group
                  | Table
                  | Function
                  | Resource
                  | PrefixOp Operand

/* Keywords and Special Identifiers */
Keyword      ::= "this" | "left" | "right"

/* Complex Structures */
Group        ::= "(" Expression ")"

/* Tables */
/* Tables are constructed using [ ] or by using the comma
operator in a Group */
Table         ::= "[" Expression* "]"

/* Functions (Lambdas / Blocks) */
/* Note: LBP/RBP Integers must be immediately adjacent to
braces (no spaces) */
Function      ::= (INTEGER)? "{" Expression "}" (INTEGER)?

/* Resources */
Resource     ::= Expression "@:" Table

```

```

/* Operators */
Operator   ::= IDENTIFIER
              | "+" | "-" | "*" | "/" | "%" | "***"
              | "=" | "<>" | "<" | ">" | "<=" | ">="
              | "&&" | "||" | "!" | "~"
              | "&" | "|" | "^" | "<<" | ">>"
              | "->" | "-<" | "-<>"
              | "." | "?" | "?:" | "??"
              | ":" | "," | "o"
              | "++" | "--" | "@"

PrefixOp   ::= IDENTIFIER | "-" | "!" | "~" | "@" | "++" |
              "--"

```

Build Model

The OrgLang compiler (`org`) operates on a single entrypoint file and resolves dependencies recursively.

Entrypoint and Output

- **Executable:** The compiler always produces an executable binary by default. There is no distinction between “library” and “executable” projects at the build level.
- **Command:** `org build main.org` compiles `main.org` (and its dependencies) into an executable named `main`.
- **Output Name:** You can specify the output filename using the `-o` or `--output` flag: `org build main.org -o myapp`.
- **Execution:** The entrypoint file is executed from top to bottom. Modules imported by it are executed when the `path @ org` expression is encountered.

Dependency Resolution

When a module is imported via `"path" @ org`:

1. **Relative Path:** The compiler first attempts to resolve the path relative to the **source file** that contains the import statement.
2. **Fallback:** If not found, it attempts to resolve relative to the current working directory or absolute paths.
3. **Compilation:** All imported modules are compiled into the single output binary. Cycles are currently allowed but will result in infinite recursion at runtime if not handled carefully (though the import cache prevents re-execution of the top-level scope).

Imports

Imports in OrgLang are expressions that evaluate to a **Table** containing the definitions from the imported file.

Import Syntax

```
lib : "path/to/lib.org" @ org;
```

Execution Semantics

- **Dynamic Evaluation:** The import is an expression evaluated at runtime.
- **No Caching:** Each time the `@ org` expression is evaluated, the module's body is executed. There is no built-in singleton cache. If a module has side effects (e.g., printing to stdout), importing it multiple times will trigger those side effects each time.
- **Optimization:** The compiler includes the module's code only once in the final binary, but the *call* to that code happens on every import evaluation.
- **Best Practice:** Assign the result of an import to a variable (`lib`) and reuse that variable to avoid re-execution.

Path Resolution

As described in the Build Model, paths are resolved relative to the source file.

Project Structure

There is no mandated project structure, but the following convention is recommended:

```
project/
├── main.org          # Entrypoint
├── lib/              # Library modules
│   ├── math.org
│   └── utils.org
├── test/             # Tests
│   └── main_test.org
└── build/            # Output directory (optional)
```

Future Goals (Wishlist)

The following features are currently **not implemented** but are planned for future versions:

1. Language Features

- **Static Analysis:** A pass to detect undefined variables and type mismatches before compilation.
- **Pattern Matching:** Enhanced syntax for destructuring Tables in function arguments.
- **Coroutines:** First-class support for suspending and resuming execution contexts.

2. Tooling

- **REPL:** An interactive Read-Eval-Print Loop for quick experimentation.
- **LSP:** A Language Server Protocol implementation for IDE support.
- **Package Manager:** A tool (`org get`) to manage external dependencies.

3. Optimizations

- **Bytecode Interpreter:** Alternatively to C transpilation, a direct bytecode interpreter for faster development cycles.
- **Tail Call Optimization:** To support deeper recursion safely.

Non-Goals (v1)

- **Strict Static Typing:** OrgLang v1 is dynamically typed. Optional type hints may be added later, but strict enforcement is not a priority.
- **Object-Oriented Classes:** The Table-based prototype system is the primary abstraction mechanism.